Module 13 Final Assignment: Applying CDC to Databases on Different Platforms

Miguel Brito

1. Creating a folder for the final assignment with the name "Assignment 13"

I'm still going to use the "Activity 13.3" folder as all the necessary logic is available there. At the end of this activity, I will add a link to GitHub so that the implementation that I used for this activity is available to be fully seen.

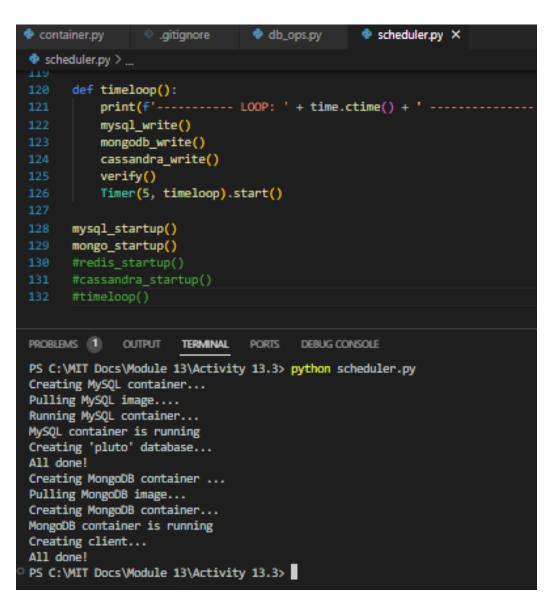
2. Displaying the logic used to create both "MySQL" and "MongoDB" containers and showing the code running to create them.

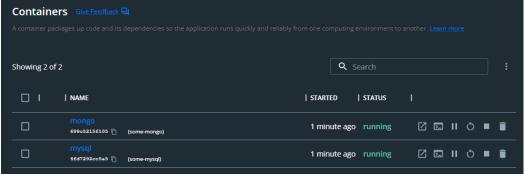
```
container.py X
                .gitignore
                                db_ops.py
                                               scheduler.py
🌵 container.py > ધ mysql
      import docker
      from docker.errors import DockerException, NotFound
     client = docker.from_env()
      class mysql:
         def __init__(self, cn_name):
             self.cn_name = cn_name
          def create(self, db_password):
             print("Pulling MySQL image.
             client.images.pull('mysql')
            print("Running MySQL container...")
             container = client.containers.run(
                 name = self.cn_name,
                 environment={"MYSQL_ROOT_PASSWORD": db_password},
                  ports={'3306/tcp': 3306},
                  detach=True
              print('MySQL container is running')
          def stop(self):
              try:
                 container = client.containers.get(self.cn_name)
                 container.stop()
             except NotFound:
                  print("Container not found")
              except DockerException as e:
                 print('An error occured when stopping the container')
          def delete(self):
              try:
                  container = client.containers.get(self.cn_name)
                 container.remove(force=True)
              except NotFound:
                 print(f"Container not found.")
              except DockerException as e:
                 print(f"An error occurred: {e}")
```

```
db_ops.py X  scheduler.py
db_ops.py > ...
      class mysql_db:
         def __init__(self, db_password):
    self.db_password = db_password
              self.connection = mysql.connector.connect(
                                                          port='3306',
                                                          database="",
                                                           auth_plugin='mysql_native_password'
              self.cursor = self.connection.cursor()
          def create(self):
              query = ("DROP DATABASE IF EXISTS 'pluto';")
              self.cursor.execute(query)
              query = ("CREATE DATABASE IF NOT EXISTS pluto")
              self.cursor.execute(query)
             query = ("USE pluto")
             self.cursor.execute(query)
              query = ('''
              CREATE TABLE posts(
                 id VARCHAR(36),
                  stamp VARCHAR(20)
              self.cursor.execute(query)
              self.connection.commit()
              self.cursor.close()
          def write(self):
             query = ("USE pluto")
              self.cursor.execute(query)
              id = str(uuid.uuid4())
             time = datetime.now().strftime('%Y-%m-%d %H:%M:%S')
              query = (f'INSERT INTO 'posts' VALUES("{id}","{time}")')
              self.cursor.execute(query)
              self.connection.commit()
              self.cursor.close()
              print(f"Data added to the MySQL database: [{id}], [{time}]")
           def read(self):
               query = ("USE pluto")
               self.cursor.execute(query)
               query = ("SELECT * FROM posts ORDER BY stamp DESC LIMIT 5;")
               self.cursor.execute(query)
               stamps = []
               for row in self.cursor.fetchall():
                   print(row)
           def close_connection(self):
               if self.connection:
                    self.connection.close()
                    print("Database connection terminated")
```

```
container.py × 🌼 .gitignore
                                                  scheduler.py
                                 db_ops.py
🌵 container.py > ધ mysql
      class mongo:
          def __init__(self, cn_name):
               self.cn_name = cn_name
          def create(self):
               print("Pulling MongoDB image...")
client.images.pull("mongo")
               print("Creating MongoDB container...")
               container = client.containers.run(
                   'mongo',
                   name = self.cn_name,
ports = {'27017/tcp': 27017},
                   detach = True
               print("MongoDB container is running")
          def stop(self):
               try:
                   container = client.containers.get(self.cn_name)
                   container.stop()
               except NotFound:
                  print("Container not found")
               except DockerException as e:
                   print("An error occured when stopping the container. ")
           def delete(self):
               try:
                   container = client.containers.get(self.cn_name)
                   container.remove(force=True)
               except NotFound:
                  print(f"Container not found.")
               except DockerException as e:
                   print(f"An error occurred: {e}")
```

```
container.py
                                                scheduler.py X
               gitignore
                                db_ops.py
scheduler.py > ...
      def mysql_startup():
          print("Creating MySQL container...")
          mysql_container = mysql('some-mysql')
          mysql_container.create('MySecretPassword')
          print("Creating 'pluto' database...")
          time.sleep(20)
          sql = mysql_db('MySecretPassword')
          sql.create()
          print("All done!")
      def mongo_startup():
          print("Creating MongoDB container ...")
          mongo_container = mongo('some-mongo')
         mongo_container.create()
          print("Creating client...")
          mng = mongo_db()
          print("All done!")
```





3. Displaying the logic used to create both "Redis" and "Cassandra" containers and showing the code running to create them.

```
container.py X 💠 db_ops.py
                                scheduler.py
🌵 container.py > 😘 mysql
109 v class cassandra:
       def __init__(self, cn_name):
              self.cn_name = cn_name
        def create(self):
             print("Pulling Cassandra image...")
             client.images.pull("cassandra")
             print("Running Cassandra container...")
             container = client.containers.run(
                  'cassandra',
                 name = self.cn_name,
                 ports = {'9042/tcp': 9042},
                 detach = True
              print("Cassandra container is running")
          def stop(self):
              try:
                  container = client.containers.get(self.cn_name)
                  container.stop()
             except NotFound:
                  print("Container not found")
              except DockerException as e:
                 print("An error occured when stopping the container. ")
          def delete(self):
              try:
                  container = client.containers.get(self.cn_name)
                  container.remove(force=True)
              except NotFound:
                 print(f"Container not found.")
              except DockerException as e:
                  print(f"An error occurred: {e}")
```

```
db_ops.py X  scheduler.py
db_ops.py > ...
          def __init__(self):
            self.cluster = Cluster(['localhost'],port=9042)
self.session = self.cluster.connect()
          def create(self):
              self.session.set_keyspace('pluto')
              self.session.execute("""
                   CREATE TABLE IF NOT EXISTS post (
id text PRIMARY KEY,
                       stamp text
               ....
               self.session.shutdown()
               self.cluster.shutdown()
          def write(self):
              self.session.execute('USE pluto')
               id = str(uuid.uuid4())
              time = datetime.now().strftime('%Y-%m-%d %H:%M:%S')
self.session.execute("""INSERT INTO post (id, stamp) VALUES (%s, %s)""", (id, time))
          self.session.shutdown()
self.cluster.shutdown()
         def read(self):
               self.session.execute('USE pluto')
               rows = self.session.execute('SELECT * FROM post')
               for row in rows:
                   print(row)
          def delete(self):
    query = f"DELETE FROM pluto"
    self.session.execute(query)
```

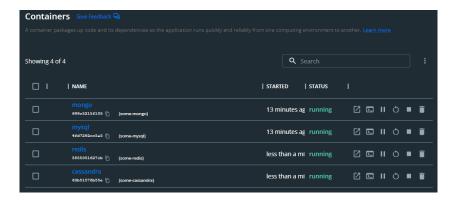
```
ontainer.py X db_ops.py
                              scheduler.py
🌵 container.py > ધ mysql
      class redis:
          def __init__(self, cn_name):
              self.cn_name = cn_name
          def create(self):
              print("Pulling Redis image")
              client.images.pull("redis")
              print("Running Redis container...")
              container = client.containers.run(
                  'redis',
                 name = self.cn_name,
                  ports = {'6379/tcp': 6379},
                  detach = True
              print("Cassandra container is running")
          def stop(self):
             try:
                  container = client.containers.get(self.cn_name)
                  container.stop()
              except NotFound:
                  print("Container not found")
              except DockerException as e:
                  print("An error occured when stopping the container.")
          def delete(self):
              try:
                  container = client.containers.get(self.cn_name)
                  container.remove(force=True)
              except NotFound:
                  print(f"Container not found.")
              except DockerException as e:
                  print(f"An error occurred: {e}")
```

```
container.py
                db_ops.py X scheduler.py
🌵 db_ops.py > ધ redis_db > 🕜 read
      class redis_db:
          def __init__(self):
               self.r = redis.Redis(host='localhost', port=6379, db=0)
          def write(self):
               id = str(uuid.uuid4())
              time = datetime.now().strftime('%Y-%m-%d %H:%M:%S')
               values = {
                   "id": id,
                   "time": time
               self.r.mset(values)
          def read(self):
               cursor = '0'
               while cursor != 0:
                   cursor, keys = self.r.scan(cursor=cursor, count=10)
                   for key in keys:
                       value = self.r.get(key)
                       print(f"{key.decode('utf-8')}: {value.decode('utf-8')}")
          def delete(self):
               keys = self.r.keys('*')
               if keys:
                   self.r.delete(*keys)
                   print("All keys deleted.")
               else:
154
                   print("No keys found to delete.")
```

```
container.py
                 db_ops.py
                                scheduler.py X
scheduler.py > ...
      def redis_startup():
          print("Creating Redis container...")
          redis_container = redis('some-redis')
          redis_container.create()
         print("Creating client...")
          rds = redis_db()
          print("All done!")
      def cassandra_startup(retries=10, delay=20):
          print("Creating Cassandra container...
          css_container = cassandra('some-cassandra')
          css_container.create()
          for attempt in range(retries):
               try:
                  print(f'Attempting to create client... (Attempt {attempt + 1} of {retries})')
                  css = cassandra_db()
                  css.create()
                  print('Cassandra setup completed successfully.')
                  break
              except NoHostAvailable:
                   print(f'Connection attempt {attempt + 1} failed. Retrying in {delay} seconds...')
                  time.sleep(delay)
          else:
              print('Failed to connect to Cassandra after multiple attempts.')
```

```
◆ container.py × ◆ db_ops.py

                                    scheduler.py X
 scheduler.py > ...
        def timeloop():
                         ------ LOOP: ' + time.ctime() + ' -----
           print(†'
            mysql_write()
            mongodb_write()
            cassandra_write()
            verify()
            Timer(5, timeloop).start()
      #mysql_startup()
      #mongo_startup()
 130 redis_startup()
 131 cassandra_startup()
 PROBLEMS (1) OUTPUT TERMINAL PORTS DEBUG CONSOLE
PS C:\MIT Docs\Module 13\Activity 13.3> python scheduler.py
 Creating Redis container...
 Pulling Redis image
Running Redis container...
 Cassandra container is running
 Creating client...
 All done!
 Creating Cassandra container...
 Pulling Cassandra image...
 Running Cassandra container...
 Cassandra container is running
 Attempting to create client... (Attempt 1 of 10)
 Connection attempt 1 failed. Retrying in 20 seconds...
 Attempting to create client... (Attempt 2 of 10)
Connection attempt 2 failed. Retrying in 20 seconds...
 Attempting to create client... (Attempt 3 of 10)
 Connection attempt 3 failed. Retrying in 20 seconds...
 Attempting to create client... (Attempt 4 of 10)
 Connection attempt 4 failed. Retrying in 20 seconds...
 Attempting to create client... (Attempt 5 of 10)
 Cassandra setup completed successfully.
O PS C:\MIT Docs\Module 13\Activity 13.3>
```



4. Providing a screenshot of the "scheduler.py" running the timeloop

I was having problems with the Timer class and had to change the way the loop works by simplifying the implementation, instead of using the "Timer" class I am using a "while True" condition as the use of this external library was provoking threading issues.

I also added a 20 second timeout in order for all the database connections to be reset and prevent errors from happening in the case of having closing connections that haven't ended their process.

However, everything still works gracefully and the registries are still added periodically.

```
PROBLEMS (1) OUTPUT TERMINAL
                                     PORTS DEBUG CONSOLE
Cassandra setup completed successfully.
 ----- LOOP: Mon Aug 12 13:02:06 2024 --
Data added to the MySQL database: [64106e6c-d7af-4ac3-a5e5-b8d99bad1592], [2024-08-12 13:02:06]
Data added to the Mongo collection: [<built-in function id>], [2024-08-12 13:02:06]
Data added to the Redis database: [2c8367f2-5c87-4ce6-893c-8cf54b301ea2], [2024-08-12 13:02:06]
Data added to the Cassandra database: [154369e1-03ea-43b3-a611-e03bede8bb09], [2024-08-12 13:02:07]
Last 5 elements added to MySQL 'pluto' database:
('64106e6c-d7af-4ac3-a5e5-b8d99bad1592', '2024-08-12 13:02:06')
Last 5 elements added to Mongo 'pluto' collection:
{'_id': ObjectId('66ba4e1e5df1da671ef31cac'), 'time': '2024-08-12 13:02:06'}
Last 5 elements added to Redis database:
time: 2024-08-12 13:02:06
id: 2c8367f2-5c87-4ce6-893c-8cf54b301ea2
Last 5 elements added to Cassandra database:
Row(id='154369e1-03ea-43b3-a611-e03bede8bb09', stamp='2024-08-12 13:02:07')
Reseting database connections...
----- LOOP: Mon Aug 12 13:02:29 2024 -----
Data added to the MySQL database: [24feaf2e-c0ee-4910-bc73-40188a8b8013], [2024-08-12 13:02:29]
Data added to the Mongo collection: [<built-in function id>], [2024-08-12 13:02:29]
Data added to the Redis database: [0ccaf35d-1430-4538-ae52-029f3c6c547c], [2024-08-12 13:02:29]
Data added to the Cassandra database: [7a050834-9827-4e37-b3c7-c802d12e699f], [2024-08-12 13:02:31]
Last 5 elements added to MySQL 'pluto' database: ('24feaf2e-c0ee-4910-bc73-40188a8b8013', '2024-08-12 13:02:29') ('64106e6c-d7af-4ac3-a5e5-b8d99bad1592', '2024-08-12 13:02:06')
Last 5 elements added to Mongo 'pluto' collection:
{'_id': ObjectId('66ba4e1e5df1da671ef31cac'), 'time': '2024-08-12 13:02:06'}
{'_id': ObjectId('66ba4e355df1da671ef31cb7'), 'time': '2024-08-12 13:02:29'}
Last 5 elements added to Redis database:
time: 2024-08-12 13:02:29
id: 0ccaf35d-1430-4538-ae52-029f3c6c547c
Last 5 elements added to Cassandra database:
Row(id='154369e1-03ea-43b3-a611-e03bede8bb09', stamp='2024-08-12 13:02:07')
Row(id='7a050834-9827-4e37-b3c7-c802d12e699f', stamp='2024-08-12 13:02:31')
Reseting database connections..
           -- LOOP: Mon Aug 12 13:02:53 2024 -
```

GitHub repo: https://github.com/MiguelBackAtItAgain/DockerContainers